Scaling & Monitoring

Cloud Computing

Brenden west

Contents

Learning Outcomes

- Elastic Load Balancing
- Amazon CloudWatch
- Amazon EC2 Auto Scaling

Reading

- AWS Cloud Foundations - Module 10

Elastic Load Balancing (ELB)

- Distributes incoming traffic across multiple targets
- Targets can be in one or more availability zones
- Targets can be EC2 instances, IP addresses, containers, or Lambda functions
-

Types of Amazon Load Balancers

Application

- Operates on application level (OSI layer 7)
- handles HTTP/HTTPs traffic
- Routes traffic based on content

Network

- Operates at network transport level (OSI layer 4)
- Works well for TCP & UDP traffic
- Can handle millions of requests per second

Classic - operates on application & transport layers. Not recommended

How Elastic Load Balancing works

- With Application Load Balancers and Network Load Balancers, you register targets in target groups, and route traffic to the target groups.
- With Classic Load Balancers, you register instances with the load balancer.

Load balancer performs health checks to monitor health of registered targets.

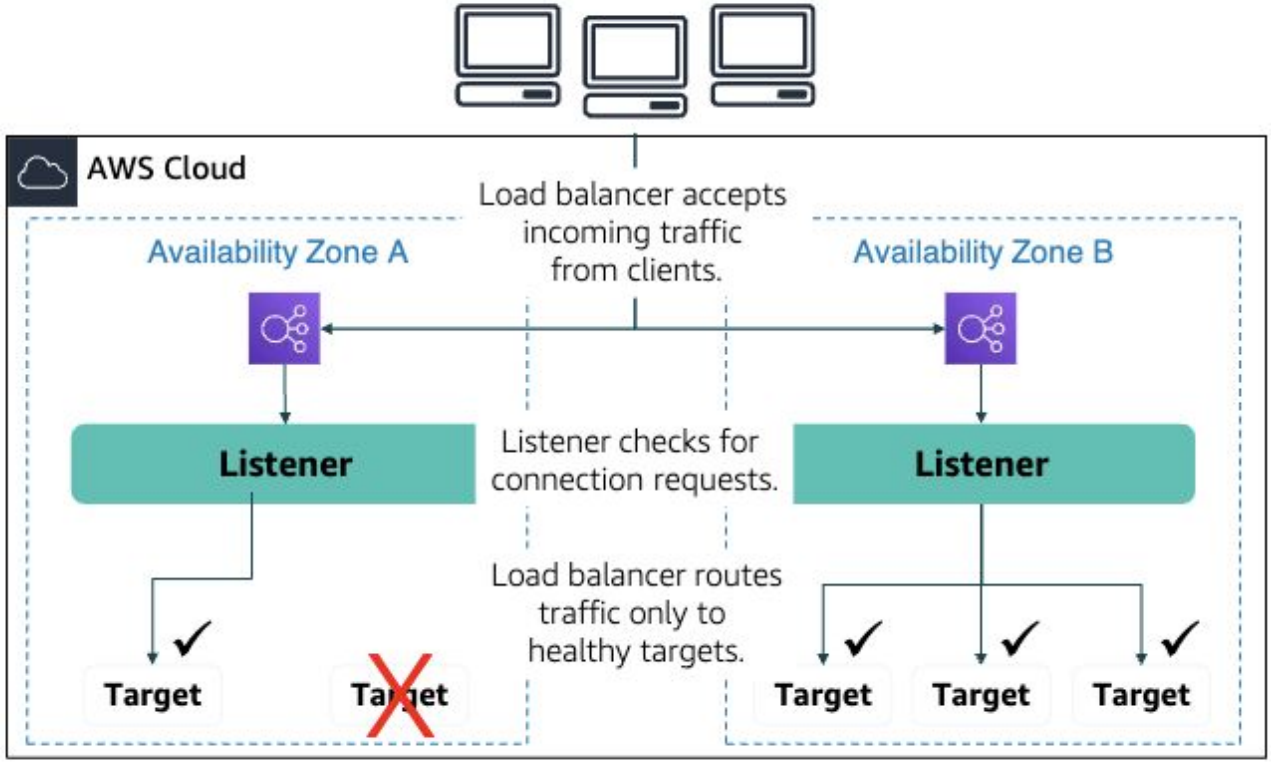

How Amazon Load Balancers work

- Configured to accept traffic with **listeners** that check for connection requests
- Registers **target groups** that can receive traffic
- Can be configured to perform health checks on registered targets.
- Only sends traffic to healthy targets
- Automatically balances load across targets
- Can register new EC2 instances provisioned by EC2 Auto Scaling
- Can be configured with security groups to control access for allowed sources
- Can invoke Lambda functions over HTTP(S)
- Integrates with CloudWatch and CloudTrail for monitoring activity

Amazon CloudWatch

- Monitors AWS resources & applications
- Collects standard & custom metrics
- Can send notifications via SNS (Simple Notification Service) or trigger EC2 Auto Scaling based on alarms
- Alarms can be based on a single metric or the result of a math expression
- Can trigger target functions based on defined events

Amazon EC2 Auto Scaling

- Automatically adds or removes EC2 instances to handle changes in demand
- Detects & replaces impaired EC2 instances & unhealthy applications
- Provides a choice of scaling options - manual, scheduled, dynamic (on-demand), & predictive
- Scaling is based on an **Auto Scaling group** with a defined policy for min, max and desired capacity
- Uses a **launch configuration** that specifies the type of EC2 instances
- ELB load balancer can be attached to the Auto Scaling group to be notified about changes to instances

Types of Scaling

- Manual -
- Scheduled - based on date & time settings
- Dynamic - based on pre-defined load parameters
- Predictive - based on predicted demand

Auto Scaling groups

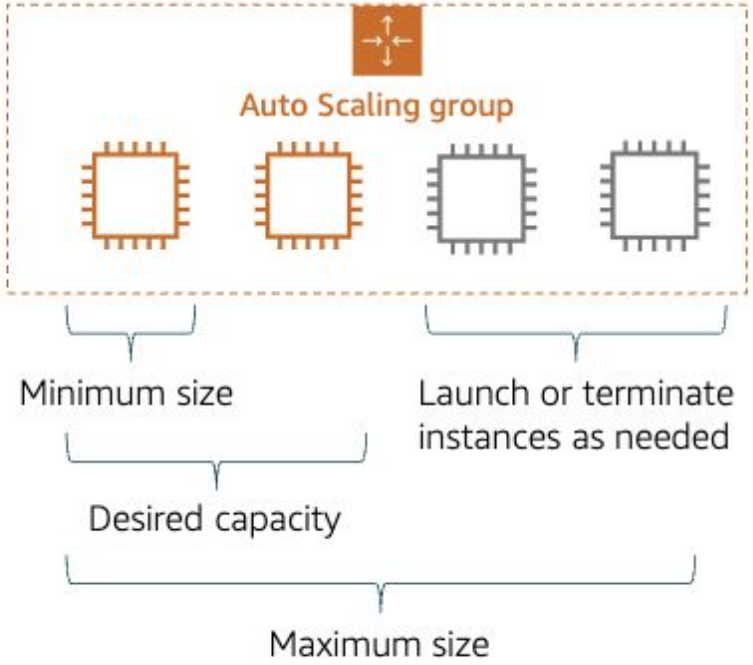

An **Auto Scaling group** is a collection of EC2 instances that are treated as a logical grouping for the purposes of automatic scaling and management.

How Amazon EC2 Auto Scaling works