
Cloud Databases

Cloud Computing
Brenden west

Contents

Learning Outcomes

- Overview of Common Datastore types
- Overview of AWS Cloud Database Offerings
- Choosing a cloud DB service
- Working with Amazon RDS
- Working with Amazon DynamoDB

Reading

- AWS Cloud Foundations - Module 8 - Databases
- AWS Cloud Developing - Module 5 - NoSQL Solutions
- <https://aws.amazon.com/free/database/>

Types of Databases

- **Relational** - Traditional DB suitable for general-purpose workloads.
 - Data stored in rows & columns.. Tables related by keys
 - Useful when strict schema & data quality required, or for transactions.
 - Hard to scale horizontally
 - E.g. - Microsoft SQL Server, Oracle, MySQL, PostgreSQL
- **Document** (aka NoSQL or schema-less) - No predefined schema structure.
 - Stores JSON **documents** in **collections**.
 - Optimal for horizontal scaling.
 - E.g. MongoDB

Relational versus non-relational databases

	Relational (SQL)	Non-Relational			
Data Storage	Rows and columns	Key-value, document, graph			
Schemas	Fixed	Dynamic			
Querying	Uses SQL	Focuses on collection of documents			
Scalability	Vertical	Horizontal			
Example					<div><pre>{ ISBN: 3111111223439, Title: "Withering Depths", Author: "Jackson, Mateo", Format: "Paperback" }</pre></div>
	ISBN	Title	Author	Format	
	3111111223439	Withering Depths	Jackson, Mateo	Paperback	
	3122222223439	Wily Willy	Wang, Xiulan	Ebook	

Other Non-relational Databases

- **In-memory** - e.g. Redis. Key-value store optimal for fast, direct access
- **Search-optimized** - e.g. Elasticsearch - Special-purpose document DB where data is indexed for complex text searches.
- **Columnar** - e.g. Redshift, AlloyDB - Column-oriented data representation stores values in a single column together. Optimized for analytical queries.
- **Graph** - data stored as a network of entities and relationships.

AWS Database Offerings

- **Relational** -
 - **Relational Database Service (RDS)** - Configurable service offering different optimizations and ~6 proprietary & open-source database engines
 - **Aurora** - Fully managed DB service compatible with MySQL and PostgreSQL
- **Document** - DynamoDB
- **Columnar** - Redshift
- **OpenSearch** - Managed search db based on Lucene engine
- **In-memory** - MemoryDB, ElastiCache - redis-compatible services
- **Graph** - Neptune

Amazon RDS

- Amazon RDS is a **managed** service that sets up and operates a relational database in the cloud.
- RDS abstracts administrative DB tasks such as server maintenance, OS & software updates & patches, DB backups
- RDS automatically scales the DB as needed
- Based on a **database instance** - an isolated DB environment that can contain multiple user-created DBs
- Pricing based on **clock hours** of running time, as well as db engine, size, & memory class

AWS RDS Instances

- Developer must choose DB engine & instance type (**on-demand** or **reserved**), backup storage, & data transfer
- Usually run on a private subnet & accessible only by authorized services
- Can be configured for **Multi-AZ deployment** with synchronous data replication
- Accessible via same tools as a standalone DB instance on EC2
- Can be configured with a **read replica** optimized for read-heavy queries

Amazon RDS DB instances

Amazon RDS



Amazon RDS DB
main instance

DB Instance Class

- CPU
- Memory
- Network performance

DB Instance Storage

- Magnetic
- General Purpose (solid state drive, or SSD)
- Provisioned IOPS

MySQL

Amazon Aurora

Microsoft SQL Server

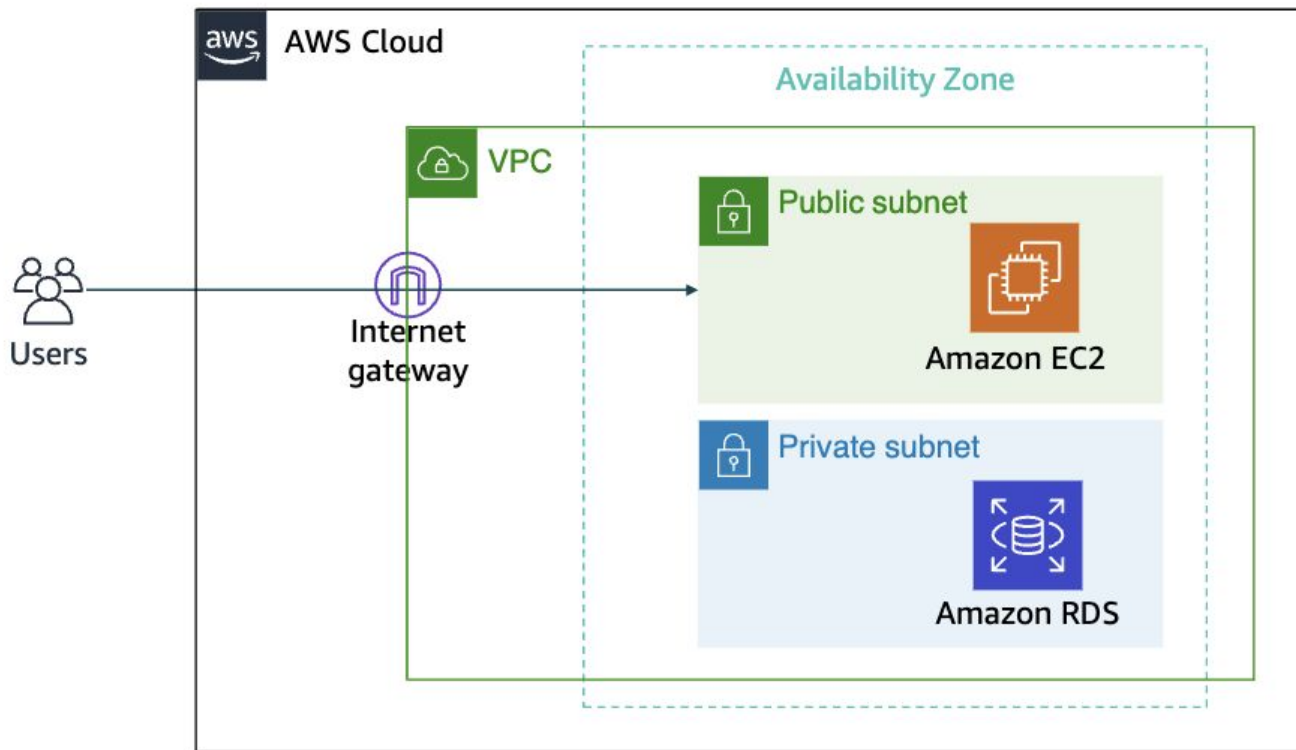
PostgreSQL

MariaDB

Oracle

DB engines

Amazon RDS in a virtual private cloud (VPC)



When to Use Amazon RDS

Use Amazon RDS when your application requires:

- Complex transactions or complex queries
- A medium to high query or write rate – Up to 30,000 IOPS (15,000 reads + 15,000 writes)
- No more than a single worker node or shard
- High durability

Do not use Amazon RDS when your application requires:

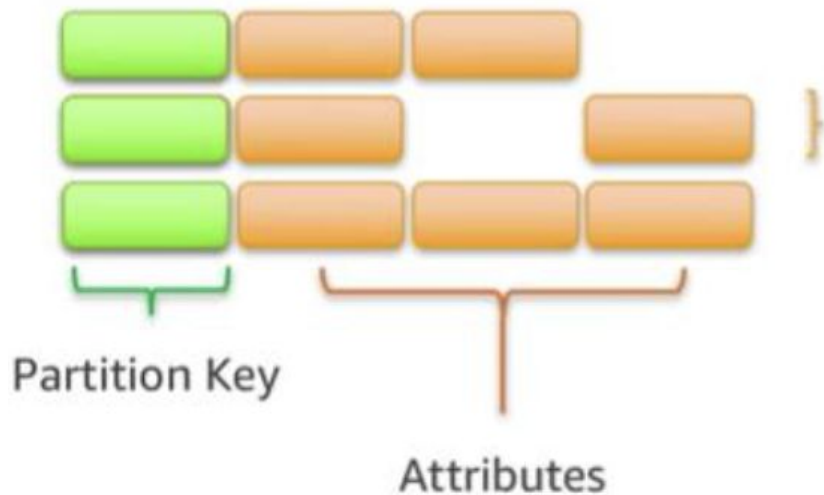
- Massive read/write rates (for example, 150,000 write/second)
- Sharding due to high data size or throughput demands
- Simple GET or PUT requests and queries that a NoSQL database can handle
- Relational database management system (RDBMS) customization

Amazon DynamoDB

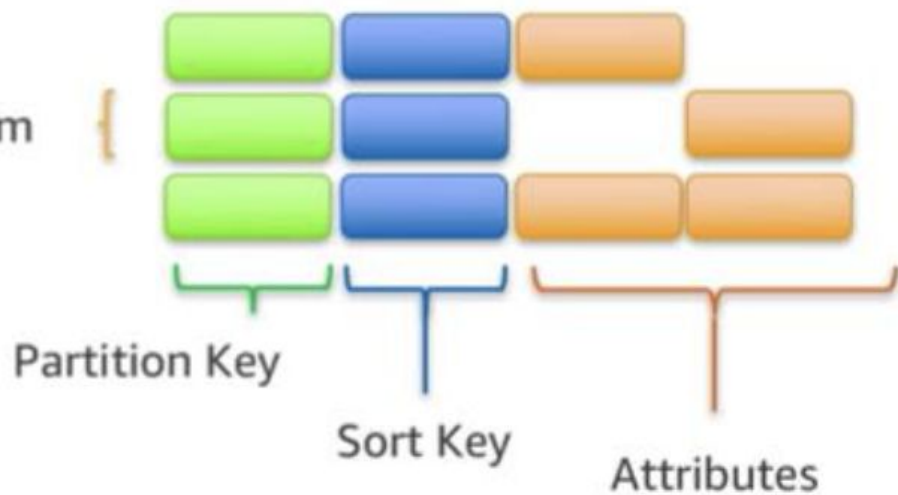
- Fast, flexible non-relational DB service w/ virtually unlimited storage
- Supports **document** & **key-value** datastore models
- Can automatically replicate across regions for redundancy
- Supports **encryption at rest** and item **time-to-live** (TTL)
- Data **items** (records) are stored in schema-less **tables** (collections)
- Supports two types of **primary key** on tables:
 - **Partition key** - single attribute **sort** key
 - **Composite primary key** - Based on two attributes
- Primary key must be a unique identifier
- As data grows, table is automatically partitioned by primary key

Items in a table must have a key

Single Key



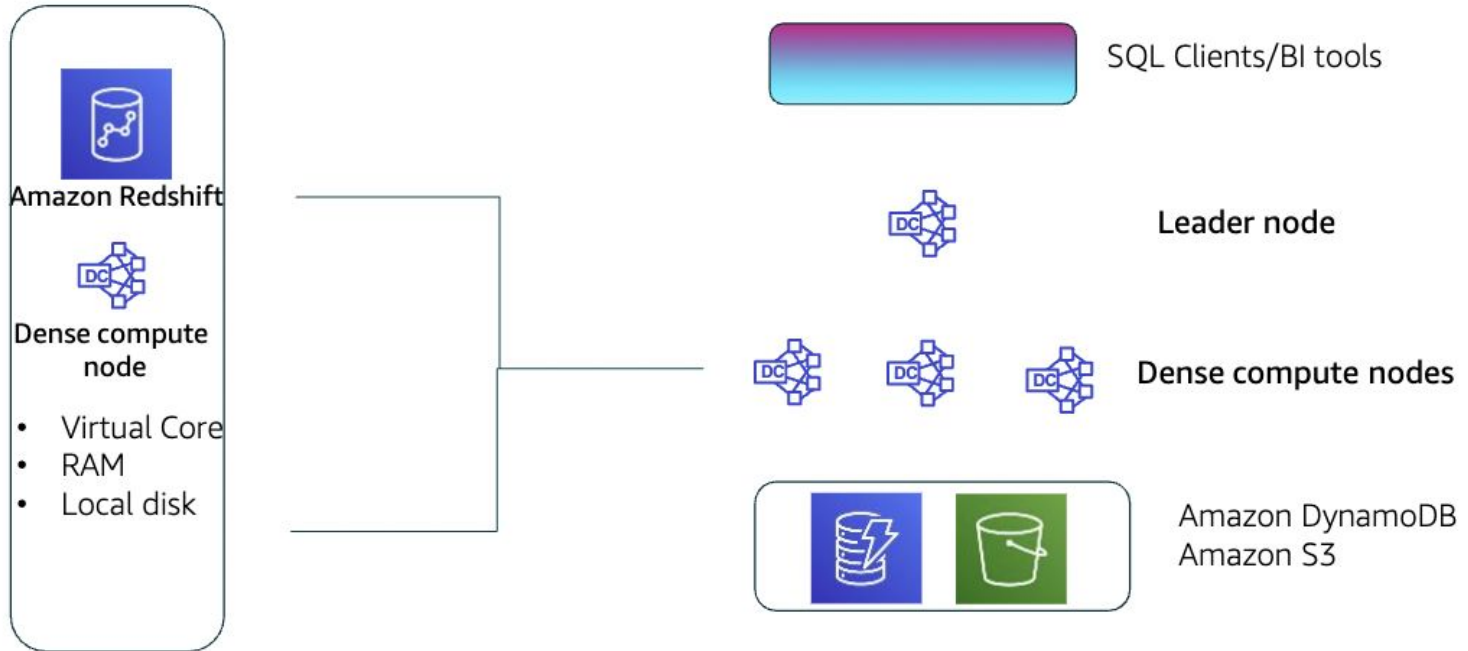
Compound Key



Amazon Redshift

- Fully managed data warehouse
- Supports standard SQL, JDBC & ODBC connectors, and common business intelligence (BI) tools
- Supports complex analytical queries using columnar storage, high-performance local disks, and parallel processing on clustered compute nodes
- Leader node manages communications with client programs and all communication with compute nodes. It develops and compiles DB query plans, assigns the code to individual compute nodes, and aggregates results

Parallel processing architecture



Amazon Aurora

- Fully managed, pay-as-you-go relational database service
- Compatible with open-source MySQL and PostgreSQL engines
- Stores multiple copies of data across multiple Availability Zones with continuous backups to AWS S3
- Designed for high performance and scalability
- Designed for instant crash recovery

The right tool for the right job

What are my requirements?

Enterprise-class relational database

Amazon RDS

Fast and flexible NoSQL database service for any scale

Amazon DynamoDB

Operating system access or application features that are not supported by AWS database services

Databases on Amazon EC2

Specific case-driven requirements (machine learning, data warehouse, graphs)

AWS purpose-built database services

Other Cloud Databases

- <https://firebase.google.com/products/realtime-database>
- <https://supabase.com/database>
- <https://www.mongodb.com/atlas/database>